

Importing libraries

```
In [4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

Reading data

```
In [5]: df = pd.read_csv('well12.csv')
df_1 = pd.read_csv('well 16A.csv')
df_2 = pd.read_excel("well F-4.xlsx", sheet_name='Drilling Parameters')
```

```
In [6]: df.head()
```

```
Out[6]:
```

	Depth (ft)	Weight on bit (k-lbs)	Hook load (k-lbs)	Surface Torque (psi)	Rotary Speed (rpm)	Flow In (gal/min)	ROP (ft/hr)
0	91.39	0.76	29.33	8.97	62.07	1102.06	450.72
1	97.75	2.13	27.96	8.97	56.05	1103.28	348.28
2	114.50	2.82	27.27	8.82	97.45	1066.39	368.77
3	117.63	0.61	29.48	8.82	89.82	1073.92	348.28
4	286.71	6.08	74.32	11.22	0.00	1037.12	58.22

```
In [7]: df_1.sample(6)
```

Out[7]:

	Time	Total Depth	Block Position	Weight on Bit	Hookload	ROP Depth/Hour	Top Drive RPM	Top Drive Torque (ft-lbs)	Flow In	Return Flow
336487	12/11/2020 18:35	8535.0	84.70821	208.38437	44.48251	0.0	0.15080	0.0	43.50702	42.35579
449167	12/29/2020 10:53	10987.0	52.04353	16.33865	202.74973	0.0	0.41245	0.0	0.00000	1.43153
479308	1/4/2021 20:57	10987.0	40.60451	186.10196	40.30299	0.0	0.00000	0.0	0.00000	0.00000
333365	12/11/2020 9:52	8535.0	84.72376	212.70111	40.17505	0.0	0.37219	0.0	0.00000	2.76115
111872	11/11/2020 14:20	5113.0	78.11756	172.26292	43.54425	0.0	0.21118	0.0	0.00000	53.22832
436017	12/27/2020 0:51	10987.0	4.07331	1.86132	36.12347	0.0	0.15080	0.0	0.00000	2.55400

First model

Merge two data sets into new data set

```
In [8]: ROP = df['ROP (ft/hr)'].append(df_1['ROP Depth/Hour'])
rpm = df['Rotary Speed (rpm)'].append(df_1['Top Drive RPM'])
WOB = df['Weight on bit (k-lbs)'].append(df_1['Weight on Bit'])
Torque = df['Surface Torque (psi)'].append(df_1['Top Drive Torque (ft-lbs)'])
```

C:\Users\IT_Center\AppData\Local\Temp\ipykernel_3252\2023585234.py:1: FutureWarning: The series.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
ROP = df['ROP (ft/hr)'].append(df_1['ROP Depth/Hour'])
```

C:\Users\IT_Center\AppData\Local\Temp\ipykernel_3252\2023585234.py:2: FutureWarning: The series.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
rpm = df['Rotary Speed (rpm)'].append(df_1['Top Drive RPM'])
```

C:\Users\IT_Center\AppData\Local\Temp\ipykernel_3252\2023585234.py:3: FutureWarning: The series.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
WOB = df['Weight on bit (k-lbs)'].append(df_1['Weight on Bit'])
```

C:\Users\IT_Center\AppData\Local\Temp\ipykernel_3252\2023585234.py:4: FutureWarning: The series.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
Torque = df['Surface Torque (psi)'].append(df_1['Top Drive Torque (ft-lbs)'])
```

```
In [6]: print(len(ROP))
print(len(rpm))
```

```
print(len(WOB))  
print(len(Torque))
```

```
491141  
491141  
491141  
491141
```

```
In [10]: data = pd.DataFrame({'ROP':ROP, 'rpm': rpm, 'WOB': WOB, 'Torque': Torque})
```

```
In [11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 491141 entries, 0 to 484031  
Data columns (total 4 columns):  
#   Column  Non-Null Count  Dtype  
---  ---  
0   ROP      491140 non-null   float64  
1   rpm      491140 non-null   float64  
2   WOB      491140 non-null   float64  
3   Torque   491140 non-null   float64  
dtypes: float64(4)  
memory usage: 18.7 MB
```

```
In [12]: data = data.dropna()  
data.sample(20)
```

Out[12]:

	ROP	rpm	WOB	Torque
221228	0.00000	0.23131	155.89081	0.0000
79453	0.00000	0.55333	149.07356	0.0000
304878	0.00000	0.19106	174.25092	0.0000
477301	0.00000	0.11055	185.18495	0.0000
462937	0.00000	0.05017	183.67418	1922.3568
67183	23.89062	73.59242	37.78608	3718.3967
20011	0.00000	0.16929	0.11667	0.0000
459589	0.00000	0.03004	183.32132	0.0000
435527	0.00000	0.41245	0.38516	0.0000
342181	116.74545	69.54699	37.58836	10228.2820
14721	0.00000	0.45113	6.67081	0.0000
82518	26.56556	39.63900	44.25723	3963.4033
15121	0.00000	0.45113	6.38984	0.0000
208752	47.23237	49.27960	45.57443	5372.6973
336298	0.00000	0.31182	208.91124	0.0000
372246	67.20520	64.95815	39.04610	11595.0550
66779	37.31055	77.03404	33.87355	4384.5720
20880	0.00000	0.14916	0.96471	0.0000
402141	0.00000	51.11111	2.19461	11979.7750
292240	0.00000	14.76264	13.24668	7616.2285

Make standard scaler to data

```
In [14]: target_1 = data['Torque']
         inputs_1 = data[['ROP', 'rpm', 'WOB']]
```

```
In [16]: scaler_standard = StandardScaler()
         input_scaled = scaler_standard.fit_transform(inputs_1)
```

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(input_scaled, target_1, test_size=0.2)
```

Train the model

```
In [18]: Reg = RandomForestRegressor()  
Reg.fit(X_train, y_train)  
Reg.score(X_test, y_test)
```


```
Out[18]: 0.8632196199836359
```

```
In [19]: y_pred = Reg.predict(X_test)
```

```
In [20]: r2 = r2_score(y_test, y_pred)  
mse = mean_squared_error(y_test, y_pred)
```

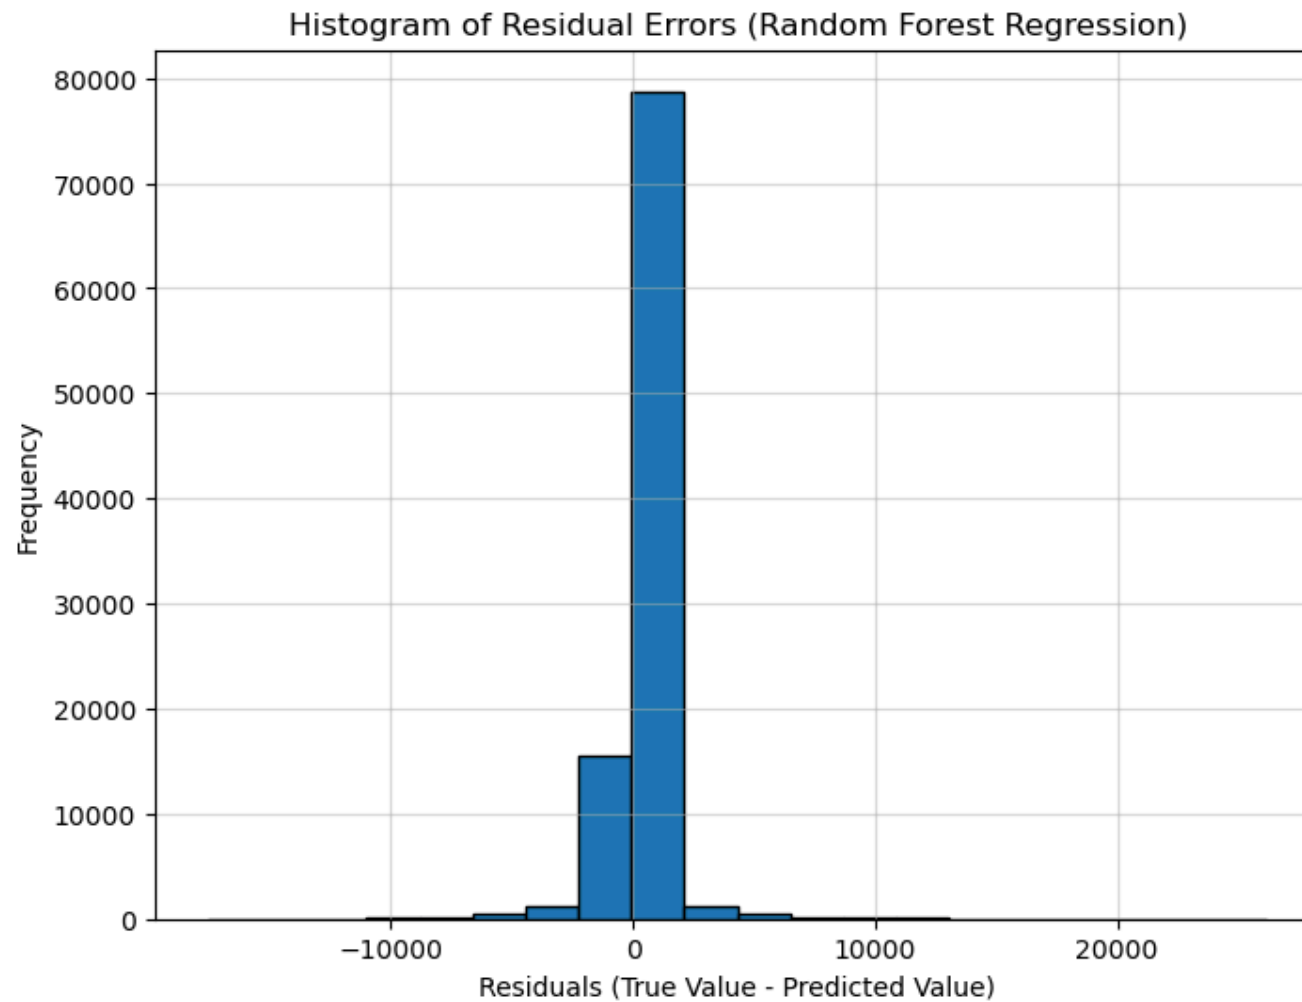
Visualize model performance

```
In [21]: plt.figure(figsize=(10, 6))  
plt.scatter(y_test, y_pred, color='blue', label='Predictions')  
plt.scatter(y_test, y_test, color='green', label='True Values', alpha=0.5)  
plt.xlabel('True Values')  
plt.ylabel('Predicted Values')  
plt.title('True Values vs. Predicted Values (Random Forest Regression)')  
plt.legend()  
plt.grid(True)  
plt.show()
```



```
In [22]: residuals = y_test - y_pred

plt.figure(figsize=(8, 6))
plt.hist(residuals, bins=20, edgecolor='black')
plt.xlabel('Residuals (True Value - Predicted Value)')
plt.ylabel('Frequency')
plt.title('Histogram of Residual Errors (Random Forest Regression)')
plt.grid(True, alpha=0.5)
plt.show()
```



Model score

```
In [25]: print(f'model score is {np.round(r2*100,2)} %', f'\nmean square error {np.round(mse,2)}') # add round
model score is 86.32 %
mean square error 1345890.43
```

Second model

Check data quality

In [26]: `df_1.isnull().sum()`

```
Out[26]: Time                0
Total Depth                1
Block Position             0
Weight on Bit              1
Hookload                   1
ROP Depth/Hour             1
Top Drive RPM              1
Top Drive Torque (ft-lbs)  1
Flow In                   1012
Return Flow                1
dtype: int64
```

In [27]: `df_1 = df_1.dropna()`

In [28]: `df_1.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 483019 entries, 0 to 484031
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Time                  483019 non-null object
 1   Total Depth           483019 non-null float64
 2   Block Position        483019 non-null float64
 3   Weight on Bit         483019 non-null float64
 4   Hookload              483019 non-null float64
 5   ROP Depth/Hour        483019 non-null float64
 6   Top Drive RPM         483019 non-null float64
 7   Top Drive Torque (ft-lbs) 483019 non-null float64
 8   Flow In               483019 non-null float64
 9   Return Flow           483019 non-null float64
dtypes: float64(9), object(1)
memory usage: 40.5+ MB
```

In [29]: `df_2.isnull().sum()`


```
Out[29]: Depth          0
         TVD            21
         ROP            21
         WOB            21
         Torque         21
         RPM            21
         SPP            21
         Flow pumps     21
         Tot Gas        21
         mW IN          21
         D Exp          21
         dtype: int64
```

```
In [30]: df_2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 675 entries, 0 to 674
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Depth           675 non-null   object
 1   TVD             654 non-null   object
 2   ROP             654 non-null   object
 3   WOB             654 non-null   object
 4   Torque          654 non-null   object
 5   RPM             654 non-null   object
 6   SPP             654 non-null   object
 7   Flow pumps      654 non-null   object
 8   Tot Gas         654 non-null   object
 9   mW IN           654 non-null   object
10   D Exp           654 non-null   object
dtypes: object(11)
memory usage: 58.1+ KB
```

```
In [31]: df_2 = df_2.dropna()
```

```
In [32]: target_2 = df_1['Hookload']
         inputs_2 = df_1[['Weight on Bit', 'ROP Depth/Hour', 'Top Drive Torque (ft-lbs)', 'Flow In', 'Block Position', 'Top Drive RPM']]
```

```
In [33]: X_train_1, X_test_1, y_train_1, y_test_1 = train_test_split(inputs_2, target_2, test_size=0.2)
```

Model training


```
In [34]: model = RandomForestRegressor()
         model.fit(X_train_1, y_train_1)
         model.score(X_test_1, y_test_1)
```

Out[34]: 0.9425561210883305

In [35]: y_pred_1 = model.predict(X_test_1)


Visualize model performance

```
In [36]: plt.figure(figsize=(10, 6))
plt.scatter(y_test_1, y_pred_1, color='blue', label='Predictions')
plt.scatter(y_test_1, y_test_1, color='green', label='True Values', alpha=0.5)
plt.xlabel('True Values')
plt.ylabel('Predicted Values')
plt.title('True Values vs. Predicted Values (Random Forest Regression)')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [37]: residuals = y_test_1 - y_pred_1

plt.figure(figsize=(8, 6))
plt.hist(residuals, bins=20, edgecolor='black')
plt.xlabel('Residuals (True Value - Predicted Value)')
plt.ylabel('Frequency')
plt.title('Histogram of Residual Errors (Random Forest Regression)')
plt.grid(True, alpha=0.5)
plt.show()
```



```
In [38]: r2_1 = r2_score(y_test_1, y_pred_1)
mse_1 = mean_squared_error(y_test_1, y_pred_1)
```

model score

```
In [40]: print(f'model score is {np.round(r2_1*100, 2)} %', f'\nmean square error {np.round(mse_1,2)}') # add round

model score is 94.26 %
mean square error 314.36
```

Third model

```
In [41]: # df_2 = df_2.iloc[1:] #uncomment to remove first row
df_2
```

```
Out[41]:
```

	Depth	TVD	ROP	WOB	Torque	RPm	SPP	Flow pumps	Tot Gas	mW IN	D Exp
22	255	255	101.54	6.63	1.38	150	86.4	3675	0	1.03	0.66
23	260	260	78.68	5.52	4.8	150	91.9	3675	0	1.03	0.7
24	265	265	45.84	2.39	3.58	150	90.5	3675	0	1.03	0.68
25	270	270	37.43	2.46	4.16	150	90.9	3675	0	1.03	0.72
26	275	275	28.63	2.47	3.4	138	77.9	3278	0	1.03	0.89
...
670	3495	3091.82	27.88	3.57	21.72	140	209.1	2227	0.28	1.4	0.74
671	3500	3094.85	27.82	3.57	21.45	140	209.9	2227	0.3	1.4	0.74
672	3505	3097.87	28.21	3.6	21.62	140	209.4	2227	0.27	1.4	0.74
673	3510	3100.9	19.33	5.9	22.39	140	209.3	2227	0.2	1.4	0.91
674	3515	3103.93	20.43	6.15	22.39	140	208.9	2227	0.15	1.4	0.9

653 rows × 11 columns

```
In [42]: df_2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 653 entries, 22 to 674
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Depth       653 non-null    object
1   TVD         653 non-null    object
2   ROP         653 non-null    object
3   WOB         653 non-null    object
4   Torque      653 non-null    object
5   RPM         653 non-null    object
6   SPP         653 non-null    object
7   Flow pumps  653 non-null    object
8   Tot Gas     653 non-null    object
9   mW IN       653 non-null    object
10  D Exp       653 non-null    object
dtypes: object(11)
memory usage: 61.2+ KB
```

convert features to appropriate data type

```
In [43]: df_2['SPP'] = df_2['SPP'].astype(float)
df_2['Flow pumps'] = df_2['Flow pumps'].astype(float)
df_2['ROP'] = df_2['ROP'].astype(float)
df_2['WOB'] = df_2['WOB'].astype(float)
```

```
In [44]: target_3 = df_2['SPP']
inputs_3 = df_2[['Flow pumps', 'ROP', 'WOB']]
```

```
In [45]: X_train_2, X_test_2, y_train_2, y_test_2 = train_test_split(inputs_3, target_3, test_size=0.2)
```

Train the model

```
In [46]: model = RandomForestRegressor()
model.fit(X_train_2, y_train_2)
model.score(X_test_2, y_test_2)
```

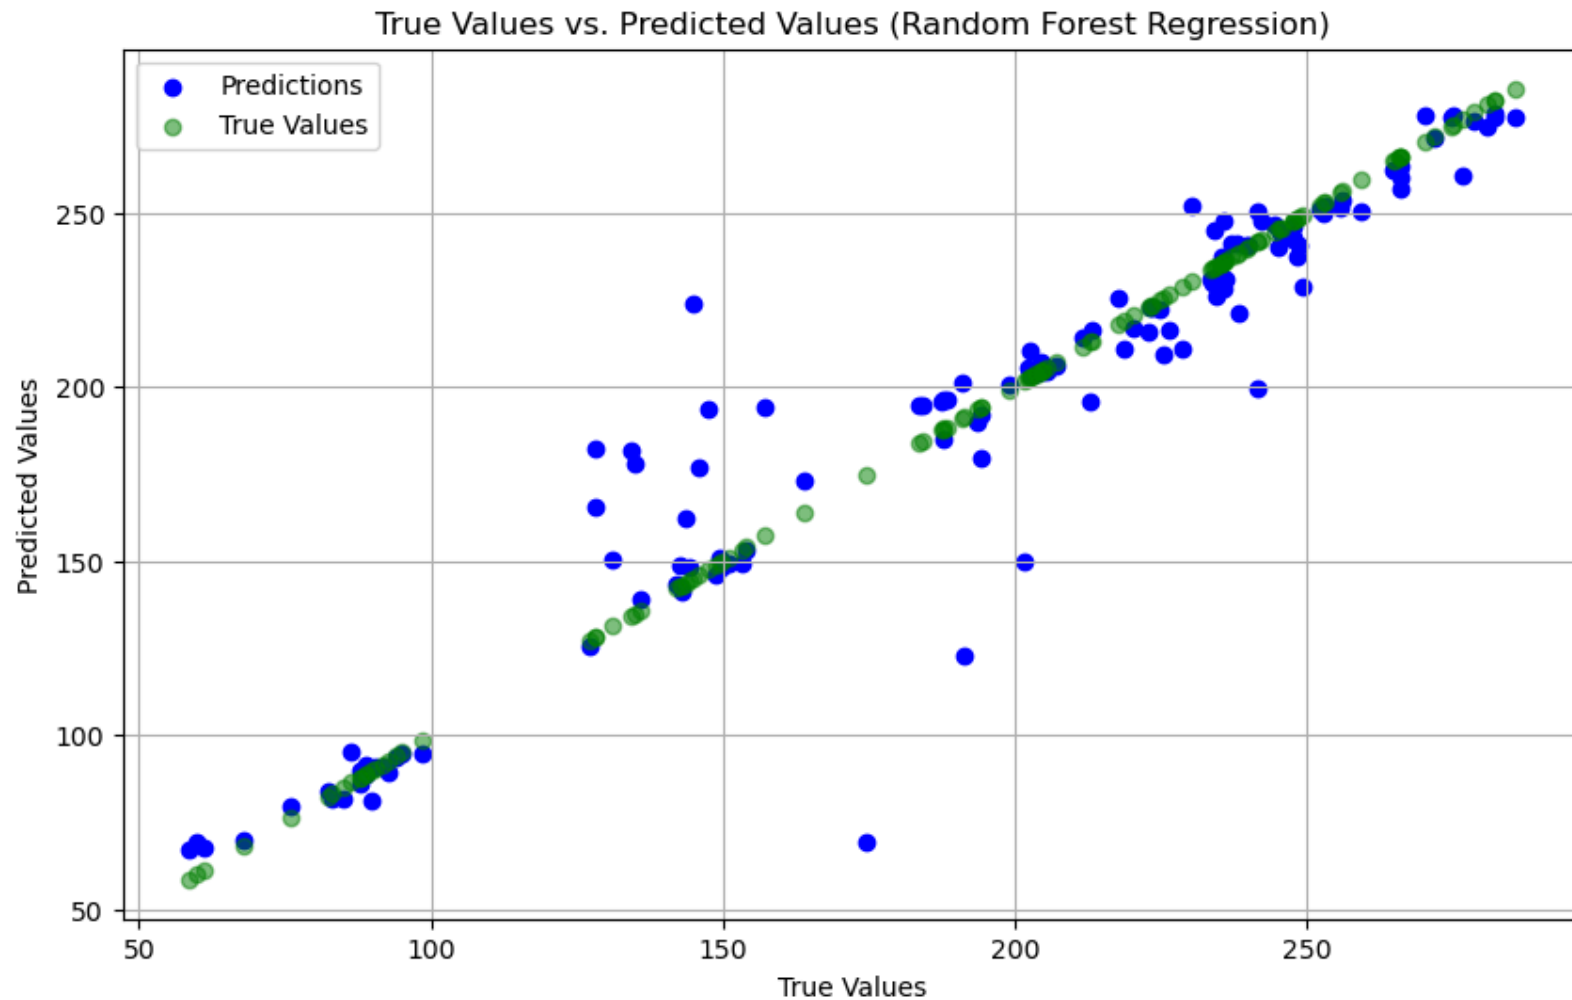
```
Out[46]: 0.9099799625057349
```

```
In [47]: y_pred_2 = model.predict(X_test_2)
```

```
In [48]: r2_2 = r2_score(y_test_2, y_pred_2)
mse_2 = mean_squared_error(y_test_2, y_pred_2)
```

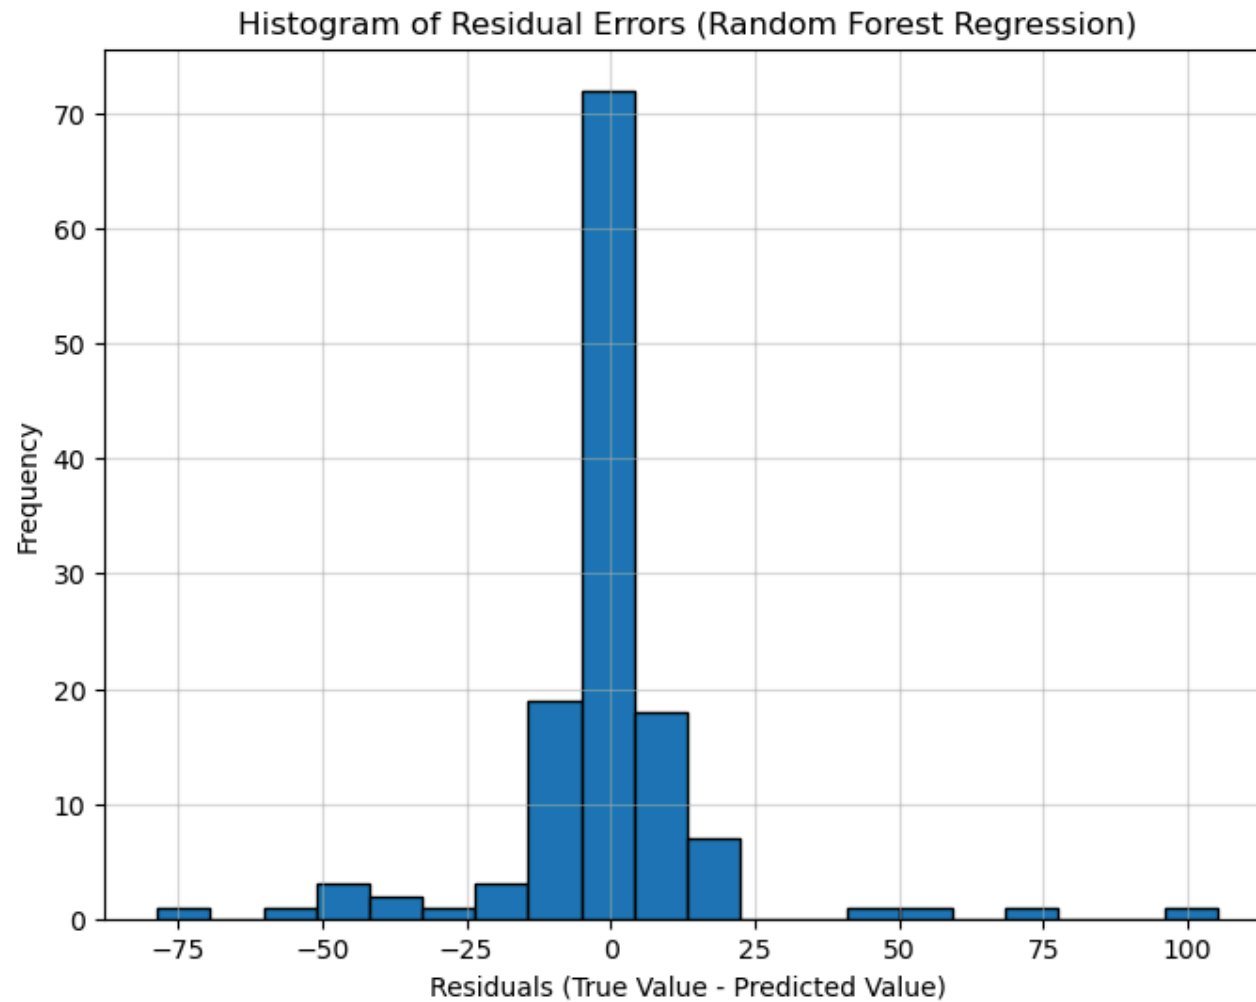
Visualize model performance

```
In [49]: plt.figure(figsize=(10, 6))
plt.scatter(y_test_2, y_pred_2, color='blue', label='Predictions')
plt.scatter(y_test_2, y_test_2, color='green', label='True Values', alpha=0.5)
plt.xlabel('True Values')
plt.ylabel('Predicted Values')
plt.title('True Values vs. Predicted Values (Random Forest Regression)')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [50]: residuals = y_test_2 - y_pred_2

plt.figure(figsize=(8, 6))
plt.hist(residuals, bins=20, edgecolor='black')
plt.xlabel('Residuals (True Value - Predicted Value)')
plt.ylabel('Frequency')
plt.title('Histogram of Residual Errors (Random Forest Regression)')
plt.grid(True, alpha=0.5)
plt.show()
```

Model score

```
In [52]: print(r2_2, mse_2)
print(f'model score is {np.round(r2_2*100, 2)} %', f'\nmean square error {np.round(mse_2, 2)}')
```

0.9099799625057349 347.14374749618304
model score is 91.0 %
mean square error 347.14

In []: